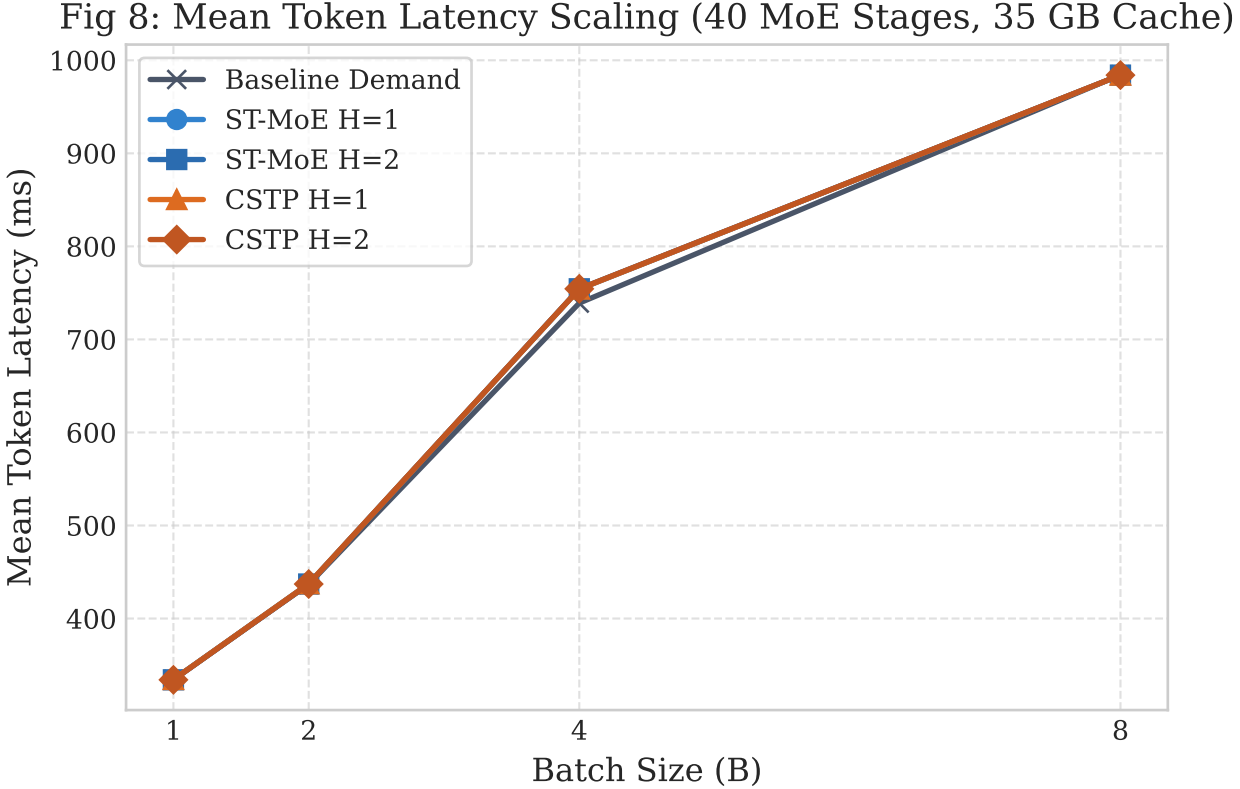
Fig 8: Mean Token Latency Scaling (40 MoE Stages, 35 GB Cache)
Baseline Demand
ST-MoE H=1
ST-MoE H=2
CSTP H=1
CSTP H=2
Mean Token Latency (ms)
1000
900
800
700
600
500
400
1
2
4
8
Batch Size (B)